STABLE

AI RELIABILITY ENGINE

STABLE

Detect wrong, risky, and unreliable AI outputs before they reach users.

Key Points

- Reliability infrastructure for production AI systems
- Highlights risky text, explains failures, and scores output confidence

STABLE

THE PROBLEM

AI is shipping faster than teams can trust it.

LLMs are moving into customer workflows, but failure modes remain invisible until they cause damage.

Key Points

- Hallucinated facts sound plausible
- Unsafe advice can reach end users
- Wrong answers erode trust, increase legal risk, and slow adoption

STABLE

WHY NOW

The market has moved from model access to model trust.

Every company is becoming an AI product company. Reliability is now the bottleneck to deployment.

Key Points

- AI is entering support, operations, healthcare, legal, and finance
- Enterprises now demand auditability and guardrails
- Teams need infrastructure, not just prompts

STABLE

OUR SOLUTION

STABLE is a reliability layer between model output and real-world use.

We generate or ingest an AI response, then run a second-pass evaluator to detect risky or unreliable content.

Key Points

- Flags hallucinations, unsafe output, weak factual grounding, and misleading certainty
- Highlights the exact text that may be wrong
- Returns confidence score, risk level, and explanation in structured form

STABLE

PRODUCT

The product makes AI failure visible in seconds.

Instead of pass or fail, STABLE shows exactly what broke and why.

STABLE

WHAT WE DETECT

Failure modes we detect

Built for the real behaviors that break trust in production AI.

Key Points

- Hallucinations
- Incorrect facts
- Unsafe outputs
- Inconsistent answers
- Low-confidence reasoning
- Weak or unverifiable claims

STABLE

TRACTION

Early traction story: strong signal, fast demos, clear buyer pain.

This is still early, but the product already demonstrates a visible and urgent category need.

Demo-ready

Investor and customer presentation quality

2-pass engine

Generation plus evaluator architecture live

Clear wedge

AI startups and enterprise copilots

Key Points

- Live interactive demo built and deployable on Vercel
- Core reliability pipeline running through OpenRouter with second-pass analysis
- Design and workflow already strong enough for pilots, investor meetings, and customer discovery

STABLE

MARKET

A growing reliability layer inside the AI stack.

As LLM usage expands, every serious deployment needs trust, visibility, and control.

Category

AI evals, observability, guardrails

Buyer

CTO, Head of AI, Product, Risk, Compliance

Why now

Trust is the bottleneck to production AI

Key Points

- TAM: global AI software and infrastructure budgets expanding rapidly
- SAM: teams deploying copilots, agents, and workflow automation with external-facing risk
- SOM: early wedge in startups and mid-market teams adopting AI into operations

STABLE

CUSTOMERS

Initial wedge: teams already shipping AI into real workflows.

If an AI response reaches a user, that product needs a trust layer.

Key Points

- AI startups and agent builders
- Enterprise copilots and support automation platforms
- Healthcare, legal, fintech, and compliance-heavy AI products

STABLE

BUSINESS MODEL

Infrastructure economics with recurring expansion paths.

STABLE can monetize as a SaaS plus usage-based reliability platform.

Key Points

- Subscription plus metered pricing per response analyzed
- Enterprise tiers for dashboards, alerts, policy customization, and audit logs
- Expansion through monitoring, eval pipelines, and deployment gating

STABLE

TEAM

Lean team, sharp problem focus, product-first execution.

The founding advantage is speed: turning a painful and visible AI problem into usable infrastructure fast.

Founder DNA

Product + engineering + execution velocity

Next hires

Reliability engineer, frontend/product design, GTM

Operating style

Fast iteration with enterprise positioning

Key Points

- Founder-led product and engineering execution
- Focus on design-forward infrastructure that demos well and solves real risk
- Next hires: full-stack engineering, AI evals, and enterprise GTM

STABLE

WHY WE WIN

We are not another model. We are the trust layer.

STABLE sits in the reliability layer of the AI stack, where urgency and value both compound.

Key Points

- Clear pain with visible ROI: reduce brand risk and ship faster
- Product value is instantly understandable in demos
- Positioned as AI observability plus QA plus reliability infrastructure

STABLE

THE ASK

We are building the system companies use to trust AI in production.

Funding accelerates product depth, enterprise readiness, and pilot customer growth.

Key Points

- Use of funds: product, engineering, reliability engine, and GTM
- Goal: become the default checkpoint before AI reaches users
- Closing line: Before AI reaches users, it should pass through STABLE